Dis 3: Sampling (Cont'd); Stata Review

> Relevant textbook chapter: 5
>
> Stata handout offered by Dr. Pac can be accessed here: Handout
>
> This handout reorganizes the official handout's information, and adds more Stata resources.

1 Sampling

- Recall from discussion 1 that a sample is a subset of data taken from the population. The action of taking this subset of data to construct your sample is called **sampling**.

- Eventually, our goal is to use the sample to draw conclusion about the population (inferential statistics), so it is important that our sample resembles the population.

- Different **sampling plans** are proposed to construct the sample, weighing the benefits of the plan against the costs:

    - **Simple random sampling**: every possible sample entry has equal chance of being selected.
    - **Stratified random sampling**: separate the population into mutually exclusive sets (i.e. strata), and then draw simple random samples from each stratum.
    - **Cluster sampling**: population is first divided into groups, and then one uses simple random sampling to select groups; all observations within the selected groups thus enter the sample.

> Exercise. Which sampling plan is used in each of the following examples?
>
> 1. Categorize all Econ 310 students based on their class standing (freshman, sophomore, junior, senior, above senior), and then randomly selects 30 students from each class.
>   Stratified random sampling
> 2. Categorize all Econ 310 students based on their class standing (freshman, sophomore, junior, senior, above senior), and then randomly select 2 out of the 5 possible groups. The groups corresponding with the class standing selected are chosen as the sample.
>   Cluster sampling
> 3. Number Econ 310 students sequentially from 1 to N. Draw 50 non-repeat random positive integers that are less than or equal to N. Select the students with the same numbers.
>   Simple random sampling

- As you can already see from the exercise, factoring in the specific steps taken when sampling, some sampling plan is expected to construct a sample that more closely resembles the population than the others.

    To formally examine how far the samples are from the population, we look at two types of errors that occur:

    1. **Sampling error**: difference between the sample and the population that exists only because the observations that happen to be included in the sample.
       $\Rightarrow$ increasing the sample size reduces this error

2. **Nonsampling errors**: more serious type of error due to samples being selected improperly.

   $\Rightarrow$ increasing the sample size will NOT reduce this type of error

   Nonsampling errors can be divided into three categories:

   (a) **Errors in data acquisition**: the data is recorded wrong (due to incorrect measurement, mistake made during transcription, human errors)

   (b) **Nonresponse errors**: responses are not obtained from certain people.

   (c) **Selection bias**: some members from the target population cannot possibly be selected to be within the sample.

Exercise. Which type of error arises from the following examples?

1. You sent out a survey to all Econ 310 students via email, but some people quickly archived your email without filling out the survey.

   Nonsampling error – nonresponse error

2. You sent out a survey to all Econ 310 students via email, but some freshmen has yet to activate their UW email account, so the survey was not delivered to them.

   Nonsampling error – selection bias

3. You randomly selected 30 Econ 310 students to have them answer your survey questions. All 30 of them responded, and you did not make any mistake in recording the data. However, your result derived from the sample is still quite different from the parameter in the population.

   Sampling error

4. You randomly selected 30 Econ 310 students to have them answer your survey questions. All 30 of them responded, but you messed up the order of items in two columns of the data recorded.

   Nonsampling error – errors in data acquisition

---

2 Stata Overview

- **What is Stata?**

  – Stata is a software that allows you to analyze data statistically. You can think about it as the advanced version of Microsoft Excel.

  – Some comparable software / programming languages out there that can do what Stata does (and maybe some more) include SAS, SPSS, R, Python, MATLAB, Julia.

      * Because Stata is currently the most popular statistical package amongst economists, it is what we'll be learning and using for this class.

      * If you go on to take Econ 400 or 410, you will continue to use Stata in that class as well.

      * The Social Science Computing Cooperative (SSCC) here at UW-Madison offers some training classes in Stata, R, and Python. If you're interested in these classes, you can find more information on this website: `https://sscc.wisc.edu/sscc_jsp/training/`

- **How to access Stata?**

  You can access Stata using either one of the two following methods:

– **Installing it onto your personal laptop (recommended)**:
  Visit UW Software Library (`software.wisc.edu`) for installation guide and license & activation key. The version of Stata to install is **Stata/SE**.

– **Logging into Winstat (i.e. a remote server; great alternative for people with Chromebook or using unsupported OS)**:
  Check out the following link for information on logging into Winstat: `https://kb.wisc.edu/sscc/using-winstat`

3   Get Started Using Stata

Let's go through the following steps to get you started on using Stata.

1. Launch Stata, and let's go through how the Stata program looks like. Specifically, identify

   - where results show up
   - where can you find the list of variables
   - where can you find more information about the variables
   - where to run your commands (via either the Command panel or the Do-file editor)

2. Before running any commands, let's set up the working directory to tell Stata which folder on your laptop should Stata read data and save graphs or results to. The easiest way to do so is to go to the menu bar, and select

   ```
   File > Change working directory...
   ```

3. Just to make sure your name is somewhere in your results, use the display command to write your name in the log. For example, assuming you happen to be Lindsey Lohan, you would type the following command:

   ```
   display "Lindsey Lohan"
   ```

4. Let's now load a set of data to do some simple statistical analysis. On your Stata problem set, we tell you exactly what command you should use to load the appropriate dataset. But for today, let's just load a sample dataset known as auto:

   ```
   sysuse auto
   ```

5. Use the describe command to determine which variable in this dataset contains "Price" and which contains "Trunk space in cubic feet":

   ```
   describe
   ```

6. Use the histogram command to graph a histogram of price and assess whether the distribution is symmetric or skewed:

   ```
   histogram price
   ```

7. Save the histogram created by clicking on the "Save" button in the graph window. Make sure you save the graph as a `.png` file.

8. Use the summarize command to calculate the mean, median, and standard deviation of trunk space:

```
summarize trunk
```

9. Did the previous command have all the information you needed? If not, let's now try the same command with the detail option:

```
summarize trunk, detail
```

10. Use the correlate command to calculate the correlation coefficient between price and trunk space:

```
correlate price trunk
```

11. Use the scatter command to graph a scatterplot of price (on the y-axis) and trunk space (on the x-axis) and assess the relationship between the two variables (note that the order of variables in the following command matters):

```
scatter price trunk
```

12. We've now finished running all practice commands. To save the output printed in the Results panel, go to

```
File > Print > Results
```

and then save your Stata output as a PDF document using your OS's printing dialogue.

4 Some More Commands For Future Use

The previous section covered all the Stata commands for the statistical operations that we have learned so far in class. There are some other Stata commands that could be helpful for your Stata problem set due around the end of the semester (December 10 at 11pm). They are listed in here for you to reference to when doing your Stata problem set.

1. Use the ci command to calculate a 95% confidence interval for price:

```
ci means price, level(95)
```

The result from running this command looks like the following:

```
. ci means price, level(95)
```

Variable	Obs	Mean	Std. err.	[95% conf. interval]	
price	74	6165.257	342.8719	5481.914	6848.6

From the result table, the last two columns under the marker [95% conf. interval] is the lower and upper bound of the 95% confidence interval. Thus, the 95% confidence interval is

$$[5481.914, \quad 6848.6]$$

2. Use the ttest command to test whether the population mean of trunk space is equal to 13 using a 10% size of test:

```
ttest trunk=13, level(90)
```

The result from running this command looks like the following:

```
. ttest trunk=13, level(90)

One-sample t test

-----------------------------------------------------------------------------
Variable |      Obs        Mean    Std. err.   Std. dev.   [90% conf. interval]
---------+-------------------------------------------------------------------
   trunk |       74    13.75676    .4972381    4.277404    12.92836    14.58515
-----------------------------------------------------------------------------
    mean = mean(trunk)                                          t =   1.5219
H0: mean = 13                                  Degrees of freedom =       73

   Ha: mean < 13                 Ha: mean != 13                 Ha: mean > 13
 Pr(T < t) = 0.9338       Pr(|T| > |t|) = 0.1323         Pr(T > t) = 0.0662
```

When testing whether the population mean of trunk space is equal to 13, the null and alternative hypotheses are

$$H_0 : \mu = 13$$
$$H_1 : \mu \neq 13$$

This means that we can look at bottom middle part of the result to perform the test:

```
. ttest trunk=13, level(90)

One-sample t test

-----------------------------------------------------------------------------
Variable |      Obs        Mean    Std. err.   Std. dev.   [90% conf. interval]
---------+-------------------------------------------------------------------
   trunk |       74    13.75676    .4972381    4.277404    12.92836    14.58515
-----------------------------------------------------------------------------
    mean = mean(trunk)                                          t =   1.5219
H0: mean = 13                                  Degrees of freedom =       73

   Ha: mean < 13                 Ha: mean != 13                 Ha: mean > 13
 Pr(T < t) = 0.9338       Pr(|T| > |t|) = 0.1323         Pr(T > t) = 0.0662
```

(The Ha means "alternative hypothesis", which is the same thing as H_1. Also, in this test, we used level(90) . The option level here means confidence level. Since we have significance level – also known as size – equals to $\alpha = 10\%$, the confidence level is $1 - \alpha = 90\%$.)

In the red box highlighted, $Pr(|T| > |t|) = 0.1323$ means that the p-value for this test is 0.1323. Given that the p-value is the lowest significance level needed to reject the null, and that our significance level (i.e. size) is $10\% = 0.10 < 0.1323 =$ the lowest significance level needed for rejection of the null, this means that we fail to reject the null at 10% size.

5

3. Use the ttest command to test whether the population mean of trunk space is greater than 13 using a 10% size of test. Note: Will this command differ from the one above? Why or why not?

The command will not differ from the one above. Now our alternative hypothesis has changed to the following:

$$H_1 : \mu > 13$$

This means that we can look at bottom right part of the result to perform the test:

```
. ttest trunk=13, level(90)

One-sample t test
```

Variable	Obs	Mean	Std. err.	Std. dev.	[90% conf. interval]
trunk	74	13.75676	.4972381	4.277404	12.92836 14.58515

```
    mean = mean(trunk)                                          t =   1.5219
H0: mean = 13                                  Degrees of freedom =       73

    Ha: mean < 13                 Ha: mean != 13                   Ha: mean > 13
 Pr(T < t) = 0.9338       Pr(|T| > |t|) = 0.1323              Pr(T > t) = 0.0662
```

In the red box highlighted, $Pr(|T| > |t|) = 0.0662$ means that the p-value for this test is 0.0662. Since our significance level (i.e. size) is 10% $= 0.10 > 0.0662 =$ the lowest significance level needed for rejection of the null, this means that we reject the null at 10% size now.

4. Use the regress command to run a regression of price on trunk space (this language means price should be the dependent variable while trunk should be the explanatory variable, so the order of variables in the following command matters):

```
regress price trunk
```

The result from running this command looks like the following:

```
. regress price trunk
```

Source	SS	df	MS			
				Number of obs	=	74
				F(1, 72)	=	7.89
Model	62747229.9	1	62747229.9	Prob > F	=	0.0064
Residual	572318166	72	7948863.42	R-squared	=	0.0988
				Adj R-squared	=	0.0863
Total	635065396	73	8699525.97	Root MSE	=	2819.4

price	Coefficient	Std. err.	t	P>|t|	[95% conf. interval]
trunk	216.7482	77.14554	2.81	0.006	62.96142 370.535
_cons	3183.504	1110.728	2.87	0.005	969.3088 5397.699

(When using the `regress` command, the variable directly following `regress` is the dependent variable y, and whatever follows afterwards are the independent variable(s) x)

From the Coefficient column, we can read the βs estimated. The _cons row is for β_0. In other words, the estimated linear line is

$$\widehat{\text{price}} = \hat{\beta}_0 + \hat{\beta}_1 \text{ trunk}$$
$$= 3183.504 + 216.7482 \text{ trunk}$$